

Get started

Open in app



HOW TO  
OPEN THE BLACK BOX OF  
AUTOML

Follow

601K Followers



You have **1** free member-only story left this month. [Sign up for Medium and get an extra one](#)

# Exploring Auto-Sklearn Models with PipelineProfiler

How to open the black box of AutoML



Jorge Piazzenti Ono Jun 3, 2020 · 4 min read ★

[Get started](#)[Open in app](#)

Image credit: [Negative Space](#)

Building machine learning pipelines is often a difficult, time-consuming and trial-and-error task. AutoML makes this process easier, by automatically selecting computational steps, tuning hyperparameters and training end-to-end models that solve a machine learning problem. Auto-Sklearn<sup>1</sup> is one of the most popular open source AutoML systems. Given a dataset, a problem type (classification or regression) and a metric score, Auto-Sklearn is able to produce ensemble pipelines that optimize the chosen metric and produce good results.

Take, for example, the *Digits* dataset, where we want to classify images of numbers. We can automatically create a pipeline that solves this problem with a few lines of code:

```
import sklearn.datasets
import autosklearn.classification
X, y = sklearn.datasets.load_digits(return_X_y=True)
automl = autosklearn.classification.AutoSklearnClassifier()
automl.fit(X, y, dataset_name='digits')
```

[Get started](#)[Open in app](#)

obscure: the data scientist does not know what steps were taken to produce the pipeline, what part of the search space was explored, and what hyperparameters were tuned.

In this blog post, we show how *PipelineProfiler* can be used to debug and explore the models produced by Auto-Sklearn. *PipelineProfiler* is a new library aimed at producing detailed visualizations of end-to-end machine learning pipelines. It is integrated with Python Notebooks (Jupyter and Google Colab), and can be used within the data science workflow.

## The PipelineProfiler Tool

Given the trained models produced by Auto-Sklearn, we can use PipelineProfiler with three lines of code:

```
import PipelineProfiler
profiler_data = PipelineProfiler.import_autosklearn(automl)
PipelineProfiler.plot_pipeline_matrix(profiler_data)
```

Get started

Open in app

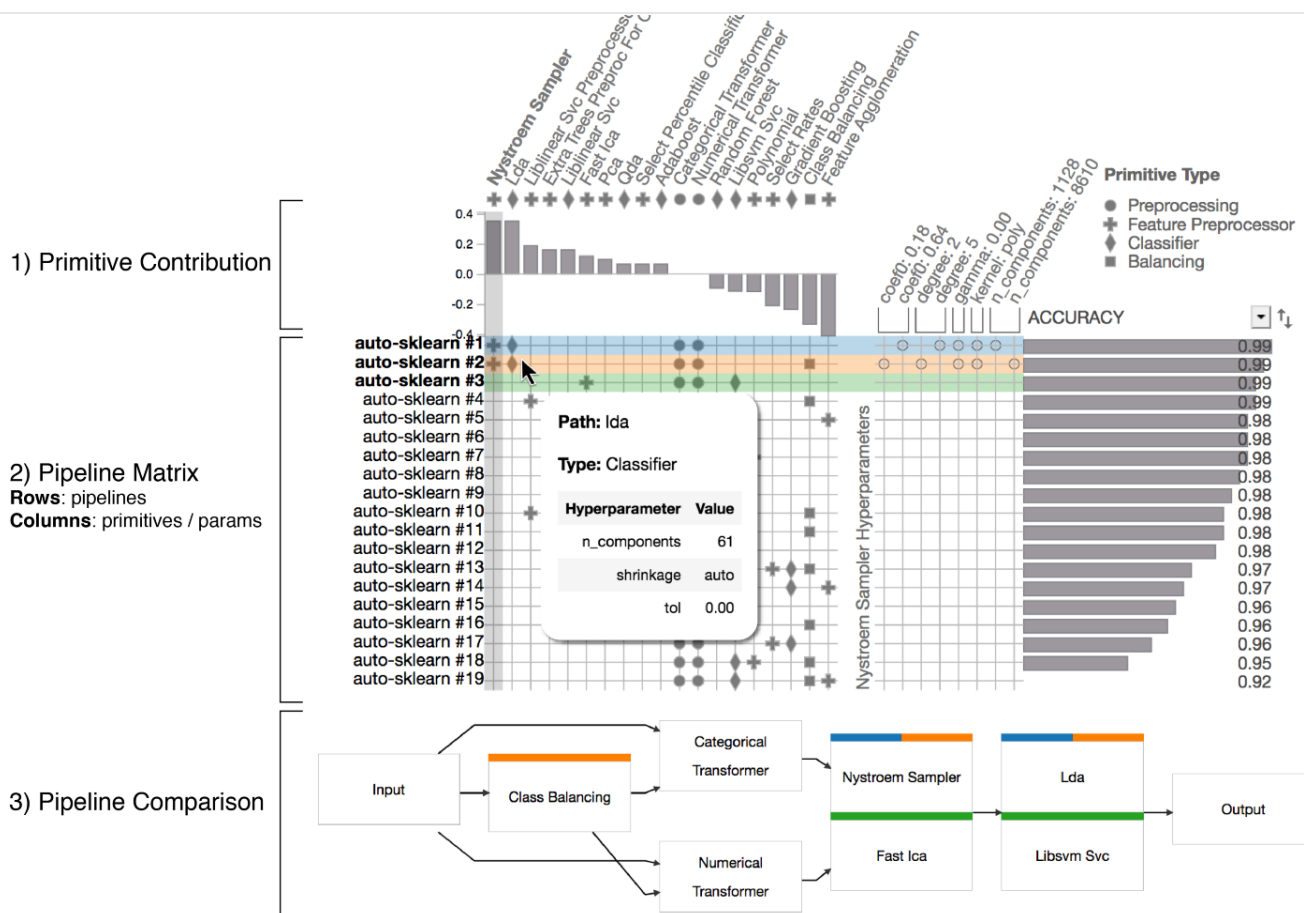


Fig.1 PipelineProfiler applied to the Digits dataset.

Fig. 1 shows *PipelineProfiler* applied to the Digits dataset. The system is divided into three parts. 1) The primitive contribution, showing the correlation of primitive usage with the pipeline scores. 2) The Pipeline Matrix, presenting a summary of the pipelines, the primitives used and the hyperparameters. 3) The Pipeline Comparison view, highlighting the differences among selected pipelines.

## Pipeline Analysis

Fig 1. shows that the top two pipelines use *Nystroem Sampler*, *LDA* and *Transformers*, but differ on the use of *Class Balancing* and the *Sampler* hyperparameters. The Primitive Contribution shows us that *Nystroem Sampler* and *LDA* are correlated with high Accuracy Scores. Conversely, *Feature Agglomeration* is correlated with low scores.

The Pipeline Comparison view presents a node-link representation of the selected pipelines (multiple pipelines can be selected by *shift-clicking*). Fig 1. shows the

[Get started](#)[Open in app](#)

*PipelineProfiler* also displays performance metrics of the evaluated pipelines. In Fig 1., we see the accuracy score of all pipelines. However, we can switch this view to other pipeline metrics. For example, the ensemble produced by Auto-Sklearn consists of a weighted average of all the evaluated pipelines. We can switch the metric score to display the ensemble weights. We can also visualize the training time of each of the pipelines. Fig 2. shows the multiple metrics that can be displayed. We removed the primitive contribution view from this image, as the contribution would be different for each of the selected metrics.

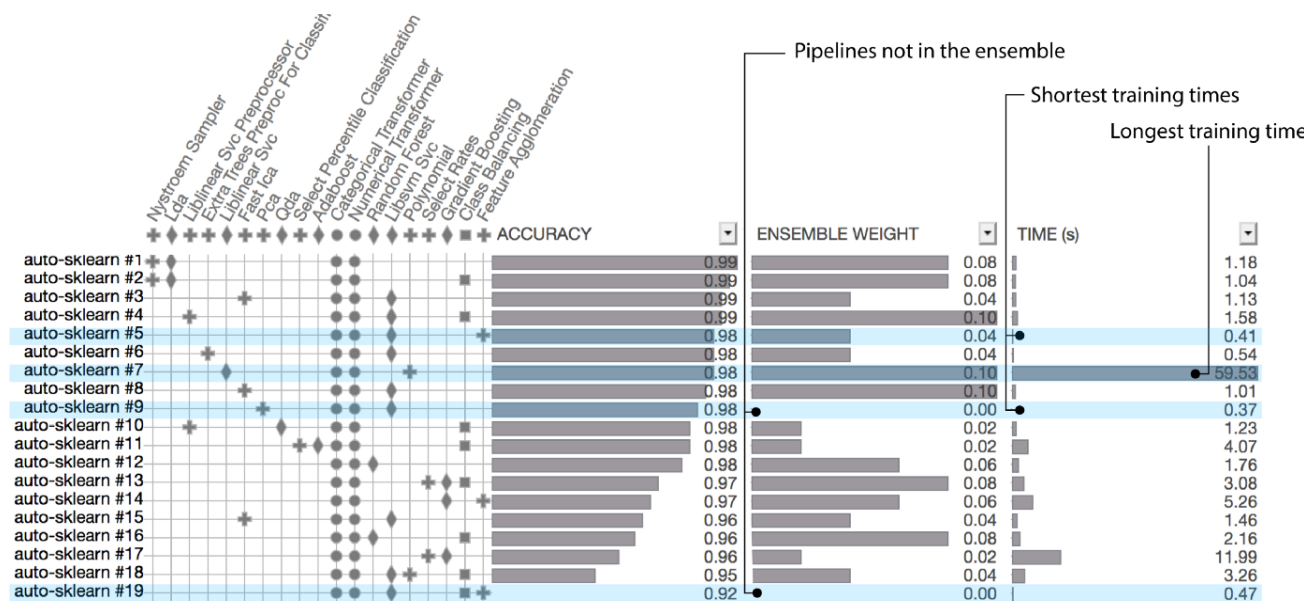


Fig. 2 PipelineProfiler can be used to inspect metrics provided by Auto-Sklearn, e.g. Accuracy, Ensemble Weight and Time.

We notice that pipelines #9 and #19 do not belong to the ensemble, as their weight is set to zero. Furthermore, we see that pipeline #7 took the longest to train, and pipelines #9 and #5 had the fastest training times. We also notice that pipeline #1 has the best accuracy, but does not have the highest ensemble weight. This shows that Auto-Sklearn uses other criteria to assign weights to pipelines in the ensemble.

If we are interested in what primitives contributed to the long training time, we can investigate the primitive contribution view:

Get started

Open in app

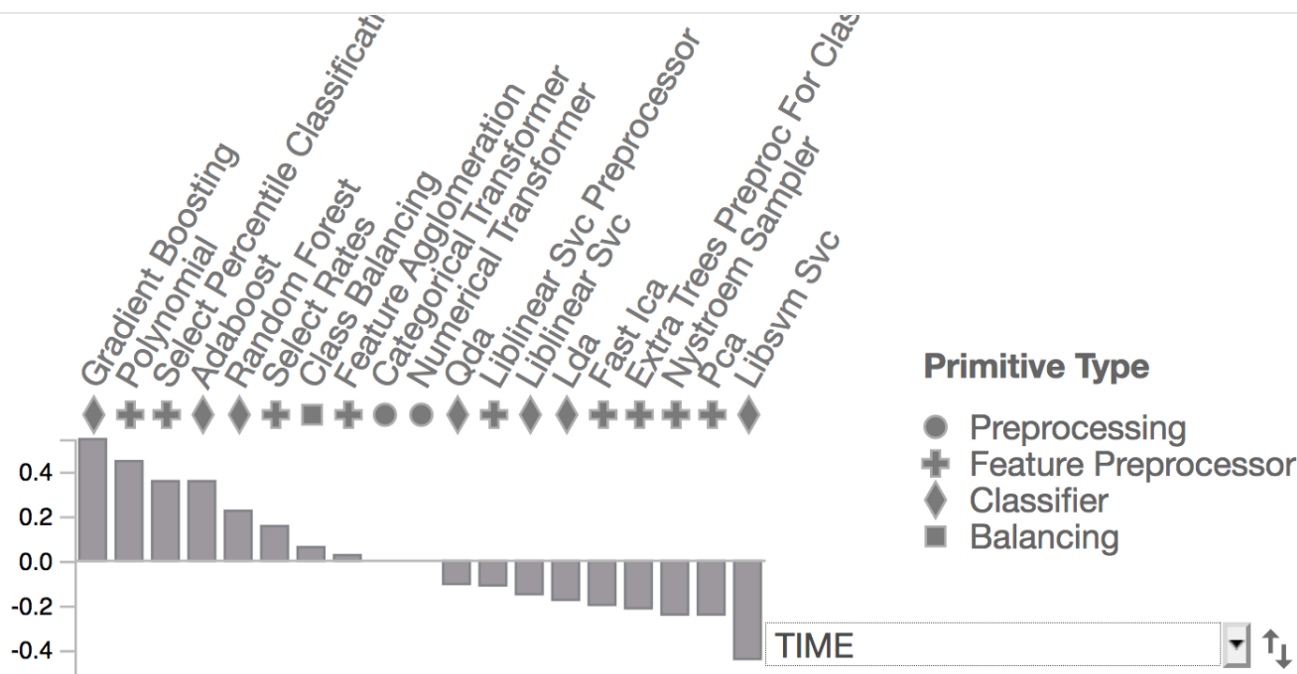


Fig 3. Primitive contribution computed based on the Time metric. Gradient Boosting is the primitive most correlated with the pipeline training time.

Here, we see that *Gradient Boosting* and *Polynomial* were two primitives correlated with the high training times, therefore they are probably to blame. Intuitively, the primitive *Polynomial* would result in longer training times, since it generates a large number of polynomial features from the original data.

## Final Thoughts

This post presented a brief introduction to *PipelineProfiler*, a Python library that can be used to explore AutoML models. We have seen how we can import and visualize pipelines from Auto-Sklearn in Jupyter Notebook.

*PipelineProfiler* is a new tool and we are actively seeking feedback on what features we should support. Let us know if you have any questions or would like to collaborate!

Get started

Open in app



4.2K views

gfyfcat

## Installation and Demo

Machine Learning | Automl | Visualization | Jupyter Notebook

The code is open sourced under the BSD 3 license. The github repository is: <https://github.com/VIDA-NYU/PipelineVis>

*PipelineProfiler* is available in the PyPI repository. It can be installed via the command:

`pip install pipelineprofiler`

Get the Medium app



at [paper](#)<sup>2</sup> and watch our video demo:



Get started

Open in app



## PipelineProfiler: A Visual Analytics Tool for the Exploration of AutoML Pipelines

